

▼ NumPy Crash Course

<https://www.youtube.com/watch?v=9JUAPgtkKpl>

<https://github.com/python-engineer/python-engineer-notebooks/blob/master/general/Numpy%20Crash%20Course.ipynb>

Learn NumPy in this complete Crash Course! I show you all the essential functions of NumPy and some tricks and useful methods.

[NumPy](#) is the core library for scientific computing in Python.

It is essential for any **data science** or **machine learning** algorithms.

Outline

1. NumPy Introduction
2. Installation and Basics
3. Array vs List
4. Dot Product
5. Speed Test array vs list
6. Multidimensional (nd) arrays
7. Indexing/Slicing/Boolean Indexing
8. Reshaping
9. Concatenation
10. Broadcasting
11. Functions and Axis
12. Datatypes
13. Copying
14. Generating arrays
15. Random numbers
16. Linear Algebra (Eigenvalues / Solving Linear Systems)
17. Loading CSV files

Installation

Installation with *pip* or *Anaconda*.

```
$ pip install numpy  
or  
$ conda install numpy
```

```
import numpy as np
```

```
np.__version__
```

```
'1.19.5'
```

Numpy is the core library in all scientific packages.

the central object in the numpy library is the numpy array.

Numpy = library for linear algebra and probability

scalar vector matrix

use cases:

- dot product/inner product
- matrix multiplication (shape must match)
- element wise matrix product
- linear systems
- inverse
- determinant
- choose random numbers (e.g. gaussian/uniform)
- working with images represented as array

Applications:

- Machine Learning
- Deep Learning
- scikit-learn
- Matplotlib
- Pandas

```
# array
a = np.array([1,2,3,4,5])
print(a)
print(a.shape) # shape of the array
print(a.dtype) # type of the elements

print(a.ndim)
print(a.size) # total number of elements
print(a.itemsize) # the size in bytes of each element
```

```
[1 2 3 4 5]
(5,)
int64
1
5
8
```

```
# essential methods
a = np.array([1,2,3])
# access and change elements
print(a[0])
a[0] = 5
print(a)
```

```
# elementwise math operations
b = a * np.array([2,0,2])
print(b)
```

```
print(a.sum())
```

```
print(a.sum())
```

```
1
[5 2 3]
[10 0 6]
10
```

```
# array vs list
```

```
l = [1,2,3]
```

```
a = np.array([1,2,3])
```

```
print(l)
```

```
print(a)
```

```
# adding new item
```

```
l.append(4)
```

```
#a.append(4) error: size of array is fixed
```

```
# there are ways to add items, but this essentially creates new arrays
```

```
l2 = l + [5]
```

```
print(l2)
```

```
a2 = a + np.array([4])
```

```
print(a2) # this is called broadcasting, adds 4 to each element
```

```
# vector addition (this is technically correct compared to broadcasting)
```

```
a3 = a + np.array([4,5,6])
```

```
print(a3)
```

```
#a3 = a + np.array([4,5]) # error, can't add vectors of different sizes
```

```
# multiplication
```

```
l2 = 2 * l # list l repeated 2 times, same as l+l
```

```
print(l2)
```

```
a3 = 2 * a # multiplication for each element
```

```
print(a3)
```

```
# modify each item in the list
```

```
l2 = []
```

```

for i in l:
    l2.append(i**2)
print(l2)
# or list comp
l2 = [i**2 for i in l]
print(l2)

a2 = a**2 # -> squares each element!
print(a2)

# Note: function applied to array usually operates element wise
a2 = np.sqrt(a) # np.exp(a), np.tanh(a)
print(a2)
a2 = np.log(a)
print(a2)

# Note: np array exists specifically to do math

[1, 2, 3]
[1 2 3]
[1, 2, 3, 4, 5]
[5 6 7]
[5 7 9]
[1, 2, 3, 4, 1, 2, 3, 4]
[2 4 6]
[1, 4, 9, 16]
[1, 4, 9, 16]
[1 4 9]
[1.          1.41421356  1.73205081]
[0.          0.69314718  1.09861229]

# dot product/inner product
a = np.array([1,2])
b = np.array([3,4])

# sum of the products of the corresponding entries
# multiply each corresponding elements and then take the sum
"
```

```
# cumbersome way for lists
dot = 0
for i in range(len(a)):
    dot += a[i] * b[i]
print(dot)
```

```
# easy with numpy :)
dot = np.dot(a,b)
print(dot)
```

```
# step by step manually
c = a * b
print(c)
d = np.sum(c)
print(d)
```

```
# most of these functions are also instance methods
dot = a.dot(b)
print(dot)
dot = (a*b).sum()
print(dot)
```

```
# in newer versions
dot = a @ b
print(dot)
```

```
11
11
[3 8]
11
11
11
11
```

```
# speed test lists vs array
from timeit import default_timer as timer

a = np.random.randn(1000)
```

```

b = np.random.randn(1000)

A = list(a)
B = list(b)

T = 1000

def dot1():
    dot = 0
    for i in range(len(A)):
        dot += A[i]*B[i]
    return dot

def dot2():
    return np.dot(a,b)

start = timer()
for t in range(T):
    dot1()
end = timer()
t1 = end-start

start = timer()
for t in range(T):
    dot2()
end = timer()
t2 = end-start

print('Time with lists:', t1)
print('Time with array:', t2)
print('Ratio', t1/t2)

Time with lists: 0.3177917980000302
Time with array: 0.001636972000028436
Ratio 194.13392409553114

# nd arrays = multidimensional arrays
# (matrix class exists but not recommended to use)

```

```
# (matrix class exists but not recommended to use)
a = np.array([[1,2], [3,4]])
print(a)
print(a.shape)

# row first, then columns
print(a[0])
print(a[0][0])
print(a[0,0])

# slicing
print(a[:,0]) # all rows in col 0
print(a[0,:]) # all columns in row 0

# transpose
print(a.T)

# matrix multiplication
b = np.array([[3, 4], [5,6]])
c = a.dot(b)
print(c)
d = a * b # elementwise multiplication
print(d)

# inner dimensions must match!
b = np.array([[1,2,3], [4,5,6]])
#c = a.dot(b.T)
#print(c)

# determinant
c = np.linalg.det(a)
print(c)

# inverse
c = np.linalg.inv(a)
print(c)

# diag
```



```
c = np.diag(a)
print(c)
```

```
# diag on a vector returns diagonal matrix (overloaded function)
c = np.diag([1,4])
print(c)
```

```
[[1 2]
 [3 4]]
(2, 2)
[[1 2]
 1
 1
 [1 3]
 [1 2]
 [[1 3]
 [2 4]]
 [[13 16]
 [29 36]]
 [[ 3  8]
 [15 24]]
-2.0000000000000004
 [[-2.  1. ]
 [ 1.5 -0.5]]
 [1 4]
 [[1 0]
 [0 4]]
```

```
# Array indexing
# Slicing: Similar to Python lists, numpy arrays can be sliced.
# Since arrays may be multidimensional, you must specify a slice for each
# dimension of the array:
a = np.array([[1,2,3,4], [5,6,7,8], [9,10,11,12]])
print(a)
```

```
# Integer array indexing
b = a[0,1]
print(b)
```

```
" a[0,1]"
```

```
# Slicing
row0 = a[0,:]
print(row0)
```

```
col0 = a[:, 0]
print(col0)
```

```
slice_a = a[0:2,1:3]
print(slice_a)
```

```
# indexing starting from the end: -1, -2 etc...
last = a[-1,-1]
print(last)
```

```
[[ 1  2  3  4]
 [ 5  6  7  8]
 [ 9 10 11 12]]
2
[1 2 3 4]
[1 5 9]
[[2 3]
 [6 7]]
12
```

```
# Boolean indexing:
a = np.array([[1,2], [3, 4], [5, 6]])
print(a)
```

```
# same shape with True or False for the condition
bool_idx = a > 2
print(bool_idx)
```

```
# note: this will be a rank 1 array!
print(a[bool_idx])
```

```
# We can do all of the above in a single concise statement:
print(a[a > 2])
```

It is not clear whether the above results are specific to the particular model used, or whether they are more general. The authors suggest that the results are likely to be general, but this needs to be confirmed by further research.

```
# np.where(): same size with modified values
```

```
b = np.where(a>2, a, -1)
```

```
print(b)
```

```
# fancy indexing: access multiple indices at once
```

```
a = np.array([10,19,30,41,50,61])
```

```
print(a)
```

```
b = a[[1,3,5]]
```

```
print(b)
```

```
even = np.argwhere(a%2==0).flatten()
```

```
print(even)
```

```
a_even = a[even]
```

```
print(a_even)
```

```
[[1 2]
```

```
 [3 4]
```

```
 [5 6]]
```

```
[[False False]
```

```
 [ True  True]
```

```
 [ True  True]]
```

```
[3 4 5 6]
```

```
[3 4 5 6]
```

```
[[-1 -1]
```

```
 [ 3  4]
```

```
 [ 5  6]]
```

```
[10 19 30 41 50 61]
```

```
[19 41 61]
```

```
[0 2 4]
```

```
[10 30 50]
```

```
# reshape
```

```
a = np.arange(1, 7)
```

```
print(a)
```

```
b = a.reshape((2, 3)) # error if shape cannot be used
```

```
print(b)
```

```
c = a.reshape((3, 2)) # 3 rows, 2 columns
```

```
print(c)
```

```
# newaxis is used to create a new axis in the data
# needed when model require the data to be shaped in a certain manner
print(a.shape)
d = a[np.newaxis, :]
print(d)
print(d.shape)
```

```
e = a[:, np.newaxis]
print(e)
print(e.shape)
```

```
[1 2 3 4 5 6]
[[1 2 3]
 [4 5 6]]
[[1 2]
 [3 4]
 [5 6]]
(6,)
[[1 2 3 4 5 6]]
(1, 6)
[[1]
 [2]
 [3]
 [4]
 [5]
 [6]]
(6, 1)
```

```
# concatenation
a = np.array([[1, 2], [3, 4]])
print(a)
b = np.array([[5, 6]])
print(b)
c = np.concatenate((a, b), axis=None)
print(c)
d = np.concatenate((a, b), axis=0)
print(d)
e = np.concatenate((a, b.T), axis=1)
```

```
c = np.concatenate((a, b), axis=1)
print(c)
```

hstack: Stack arrays in sequence horizontally (column wise). needs a tuple

```
a = np.array([1,2,3,4])
b = np.array([5,6,7,8])
c = np.hstack((a,b))
print(c)
a = np.array([[1,2], [3,4]])
b = np.array([[5,6], [7,8]])
c = np.hstack((a,b))
print(c)
```

vstack: Stack arrays in sequence vertically (row wise). needs a tuple

```
a = np.array([1,2,3,4])
b = np.array([5,6,7,8])
c = np.vstack((a,b))
print(c)
a = np.array([[1,2], [3,4]])
b = np.array([[5,6], [7,8]])
c = np.vstack((a,b))
print(c)
```

```
[[1 2]
 [3 4]]
[[5 6]]
[1 2 3 4 5 6]
[[1 2]
 [3 4]
 [5 6]]
[[1 2 5]
 [3 4 6]]
[1 2 3 4 5 6 7 8]
[[1 2 5 6]
 [3 4 7 8]]
[[1 2 3 4]
 [5 6 7 8]]
[[1 2]
 [3 4]]
```

```
[5 6]
[7 8]]
```

```
# broadcasting
```

```
# Broadcasting is a powerful mechanism that allows numpy to work with arrays of
# different shapes when performing arithmetic operations.
```

```
x = np.array([[1,2,3], [4,5,6], [7,8,9], [10, 11, 12]])
```

```
v = np.array([1, 0, 1])
```

```
y = x + v # Add v to each row of x using broadcasting
```

```
print(y)
```

```
[[ 2  2  4]
 [ 5  5  7]
 [ 8  8 10]
 [11 11 13]]
```

```
# data science functions and axes
```

```
a = np.array([[7,8,9,10,11,12,13], [17,18,19,20,21,22,23]])
```

```
print(a)
```

```
print(a.sum(axis=None)) # overall sum
```

```
print(a.sum())
```

```
print(a.sum(axis=0)) # along the rows -> 1 sum entry for each column
```

```
print(a.sum(axis=1)) # along the columns -> 1 sum entry for each row
```

```
print(a.mean(axis=None)) # overall mean
```

```
print(a.mean())
```

```
print(a.mean(axis=0)) # along the rows -> 1 mean entry for each column
```

```
print(a.mean(axis=1)) # along the columns -> 1 mean entry for each row
```

```
# some more: std, var, min, max
```

```
[[ 7  8  9 10 11 12 13]
 [17 18 19 20 21 22 23]]
210
210
[24 26 28 30 32 34 36]
[ 70 140]
15.0
15.0
```

```
[12. 13. 14. 15. 16. 17. 18.]  
[10. 20.]
```

```
# datatypes  
# https://numpy.org/devdocs/user/basics.types.html  
# Let numpy choose the datatype  
x = np.array([1, 2])  
print(x.dtype)
```

```
# Let numpy choose the datatype  
x = np.array([1.0, 2.0])  
print(x.dtype)
```

```
# Force a particular datatype, how many bits (how precise)  
x = np.array([1, 2], dtype=np.int64) # 8 bytes  
print(x.dtype)
```

```
x = np.array([1, 2], dtype=np.float32) # 4 bytes  
print(x.dtype)
```

```
int64  
float64  
int64  
float32
```

```
# copy  
a = np.array([1,2,3])  
b = a # only copies reference!  
b[0] = 42  
print(a)
```

```
a = np.array([1,2,3])  
b = a.copy() # actual copy!  
b[0] = 42  
print(a)
```

```
[42  2  3]
[1  2  3]
```

```
# generating data
a = np.zeros((2,3)) # size as tuple
b = np.ones((2,3))
print(a)
print(b)
# specific value
c = 5 * np.ones((3,3))
print(c)
c = np.full((3,3),5.0)
print(c)
# identity
d = np.eye(3) #3x3
print(d)

# arange
e = np.arange(10)
print(e)

# linspace
f = np.linspace(0, 10, 5)
print(f)
```

```
[[0.  0.  0.]
 [0.  0.  0.]]
[[1.  1.  1.]
 [1.  1.  1.]]
[[5.  5.  5.]
 [5.  5.  5.]
 [5.  5.  5.]]
[[5.  5.  5.]
 [5.  5.  5.]
 [5.  5.  5.]]
[[1.  0.  0.]
 [0.  1.  0.]
 [0.  0.  1.]]
```



```
[0 1 2 3 4 5 6 7 8 9]
[ 0.  2.5  5.  7.5 10. ]
```

```
# random numbers
a = np.random.random((3,2)) # uniform 0-1 distribution
print(a)
b = np.random.randn(3,2) # normal/Gaussian distribution, mean 0 and unit variance
# no tuple as shape here! each dimension one argument
print(b)

R = np.random.randn(10000)
print(R.mean(), R.var(), R.std())

R = np.random.randn(10, 3)
print(R.mean()) # mean of whole array

# random integer, low,high,size; high is exclusive
R = np.random.randint(3,10,size=(3,3)) # if we only pass one parameter, then from 0-x
print(R)

# with integer is between 0 up to integer exclusive
c = np.random.choice(7, size=10)
print(c)
# with an array it draws random values from this array
d = np.random.choice([1,2,3,4], size=8)
print(d)

[[0.25797747 0.71255083]
 [0.316169   0.08961938]
 [0.06743722 0.51797696]]
[[-0.43422119 -1.80315952]
 [-1.24003807 -2.28165497]
 [ 0.8919908  -0.6785463 ]]
0.011368774563225052 0.992371465236262 0.9961784304211079
-0.24087934683233125
[[6 7 8]
 [9 7 9]]
```

```
[3 8 6]]
[2 6 5 1 5 1 3 3 2 4]
[2 4 2 4 4 4 4 2]
```

```
# eigenvalues
a = np.array([[1,2], [3,4]])
eigenvalues, eigenvectors = np.linalg.eig(a)
# Note: use eigh if your matrix is symmetric (faster)
print(eigenvalues)
print(eigenvectors) # column vectors
print(eigenvectors[:,0]) # column 0 corresponding to eigenvalue[0]
```

```
# verify: e-vec * e-val = A * e-vec
d = eigenvectors[:,0] * eigenvalues[0]
e = a @ eigenvectors[:, 0]
print(d, e)
print(d == e) # numerical issues
```

```
# correct way to compare matrix
print(np.allclose(d,e))
```

```
[-0.37228132  5.37228132]
[[-0.82456484 -0.41597356]
 [ 0.56576746 -0.90937671]]
[-0.82456484  0.56576746]
[ 0.30697009 -0.21062466] [ 0.30697009 -0.21062466]
[ True False]
True
```

```
# solve linear system
# x1+x2=2200
# 1.5 x1 + 4 x2 = 5050
# 2 equations and 2 unknowns
A = np.array([[1, 1], [1.5, 4]])
b = np.array([2200,5050])
```

```
# Ax = b <=> x = A-1 b
```

```
# But: inverse is slow and less accurate
x = np.linalg.inv(A).dot(b) # not recommended
print(x)
x = np.linalg.solve(A,b) # good
print(x)
```

```
[1500.  700.]
[1500.  700.]
```

